



SAPIENZA
UNIVERSITÀ DI ROMA

Homework 2

Image Classification on Small Datasets

Machine Learning course 2021-2022



Recap

- Ensembles of small-scale convolutional networks provide good accuracy in small-data settings (Seminar 1)
- Tuning the hyper-parameters like learning rate and weight decay is a critical factor to ensure the highest performance (Seminar 2)

(Some) open questions

- Do custom/newer architectures work better than popular ResNets?
- Can we do even better choosing hyper-parameters?
- Do heterogeneous ensembles perform better than homogeneous ones?



Answering these questions with crowd-sourcing

What if each student will train a “small” network to provide an ensemble member?

In this way:


- Obtaining a (very!) large and heterogeneous ensemble
- Exploring its own search space is easier for each single member (i.e. hyper-parameters)





Dataset [\[link\]](#)

- The dataset is ciFAIR-10 of the DEIC Benchmark (Seminar 2)
 - 10-class problem
 - 50 training images per class
 - 1000 testing images per class
 - image dimension 32x32x3
- Categories:
 - airplane
 - automobile
 - bird
 - cat
 - deer
 - dog
 - frog
 - horse
 - boat
 - truck






Objectives

- Train your favorite network architecture:
 - published one (e.g. ResNet, Wide ResNet, EfficientNet...)
 - custom one (note that published ones are a good starting point)
- Analyze results concerning few modifications of chosen architecture:
 - size (e.g., network depth/width)
 - layer configurations (e.g., kernel size, convolution stride)
- Analyze results concerning few modifications of optimization process:
 - optimizer (e.g. SGD, Adam, ...)
 - optimizer parameters (e.g. learning rate, weight decay, momentum)
 - training (e.g. epochs, batch size)



Code Repository [[link](#)]

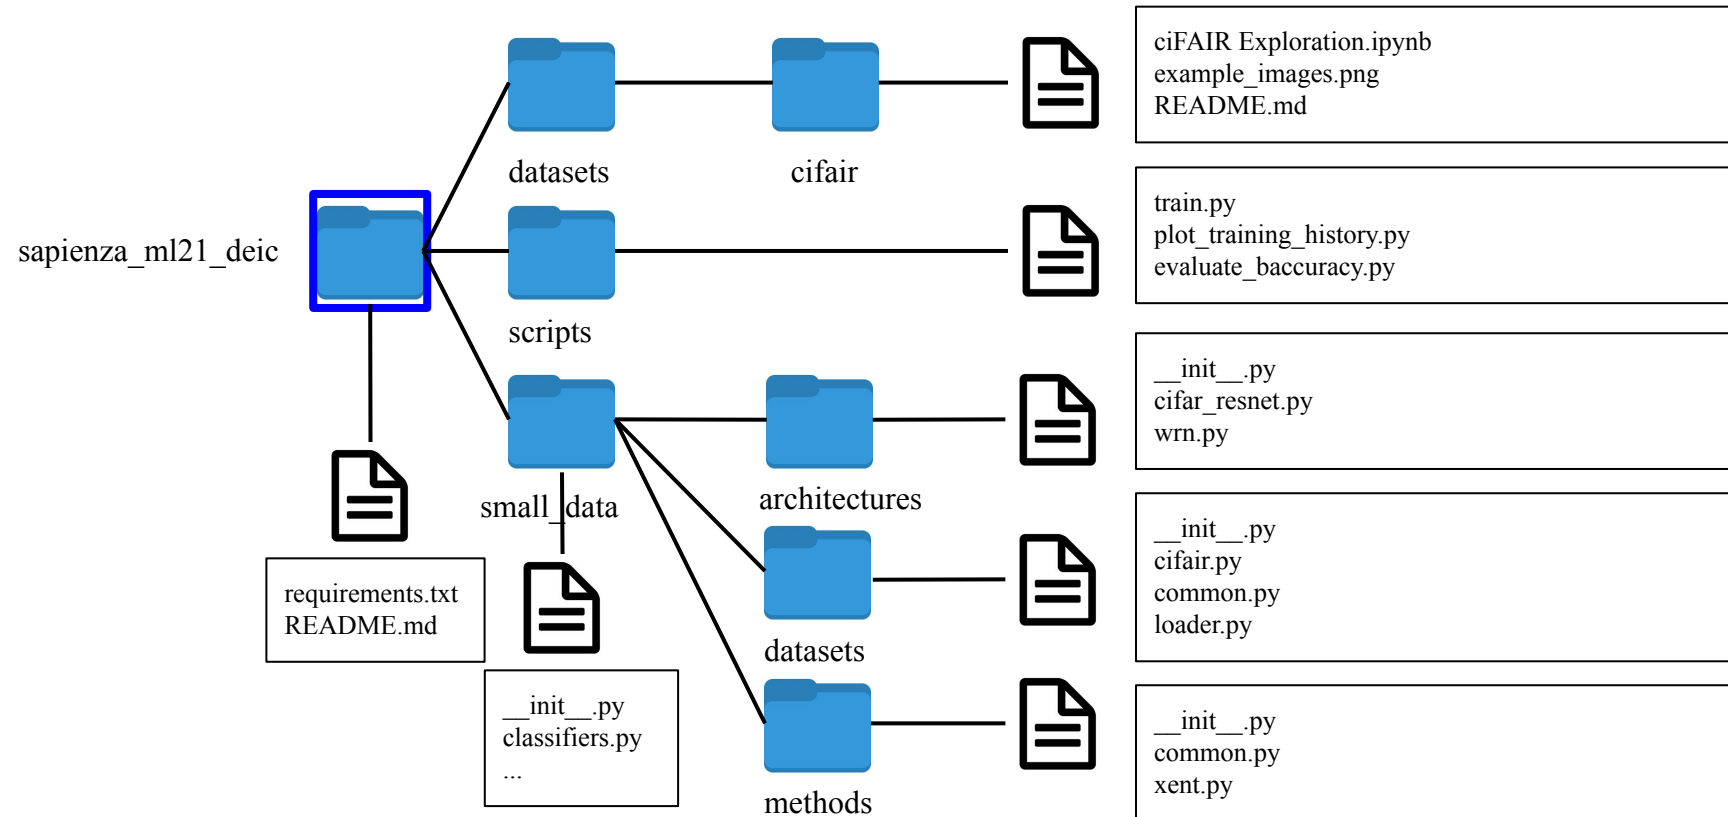
We provide a code repository as starting point for your experiments





Code Repository

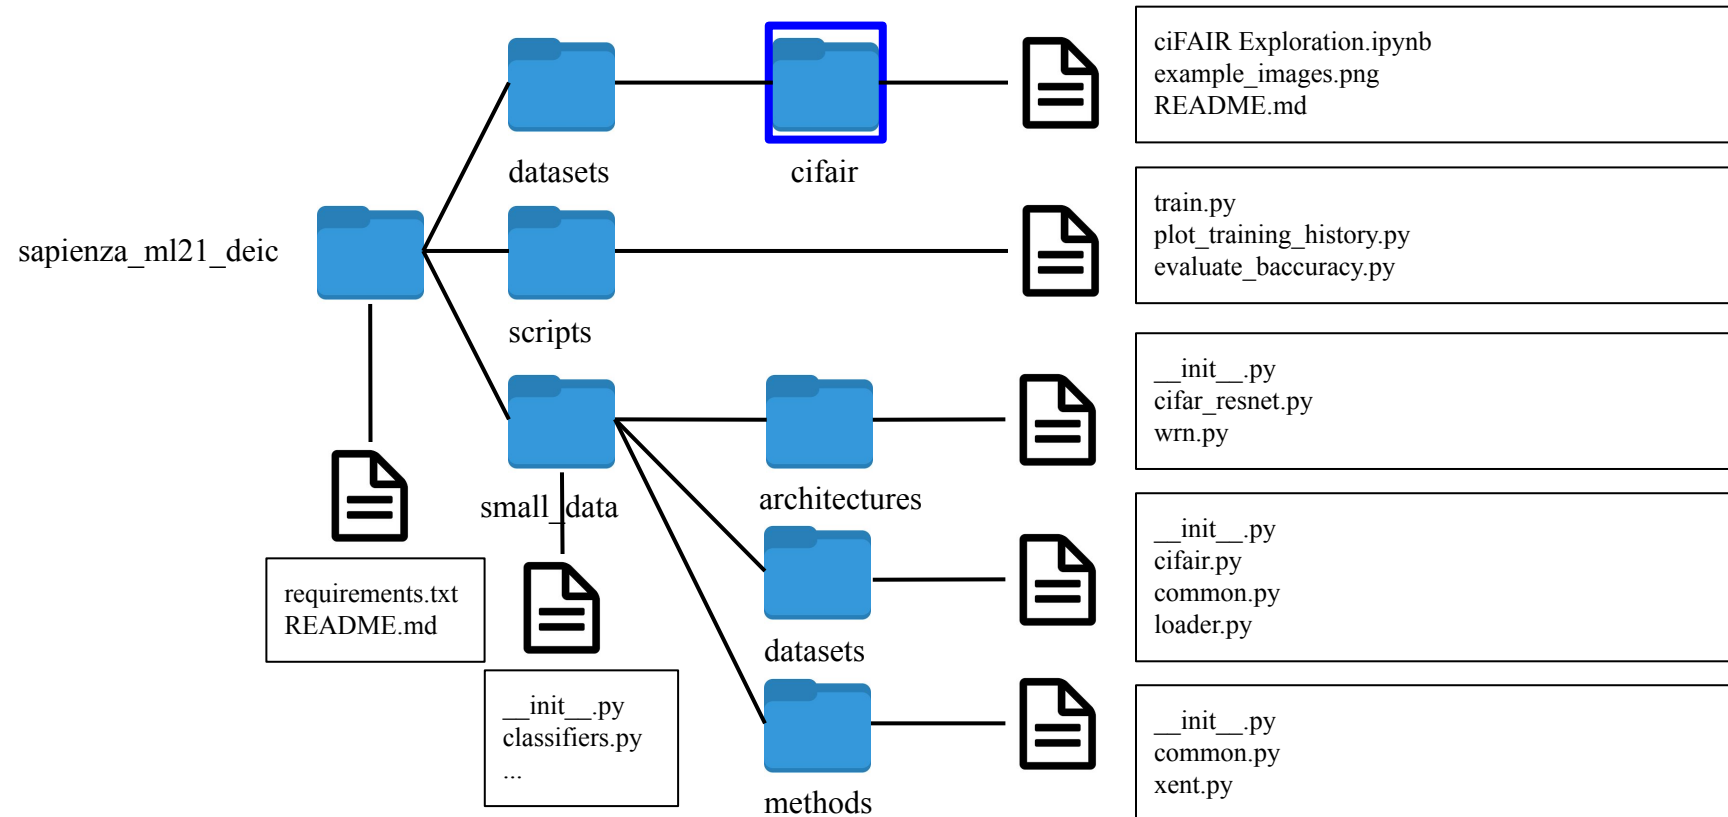
sapienza_ml21_deic is the main directory containing: file with package requirements, readme of the repository and subdirectories.





Code Repository

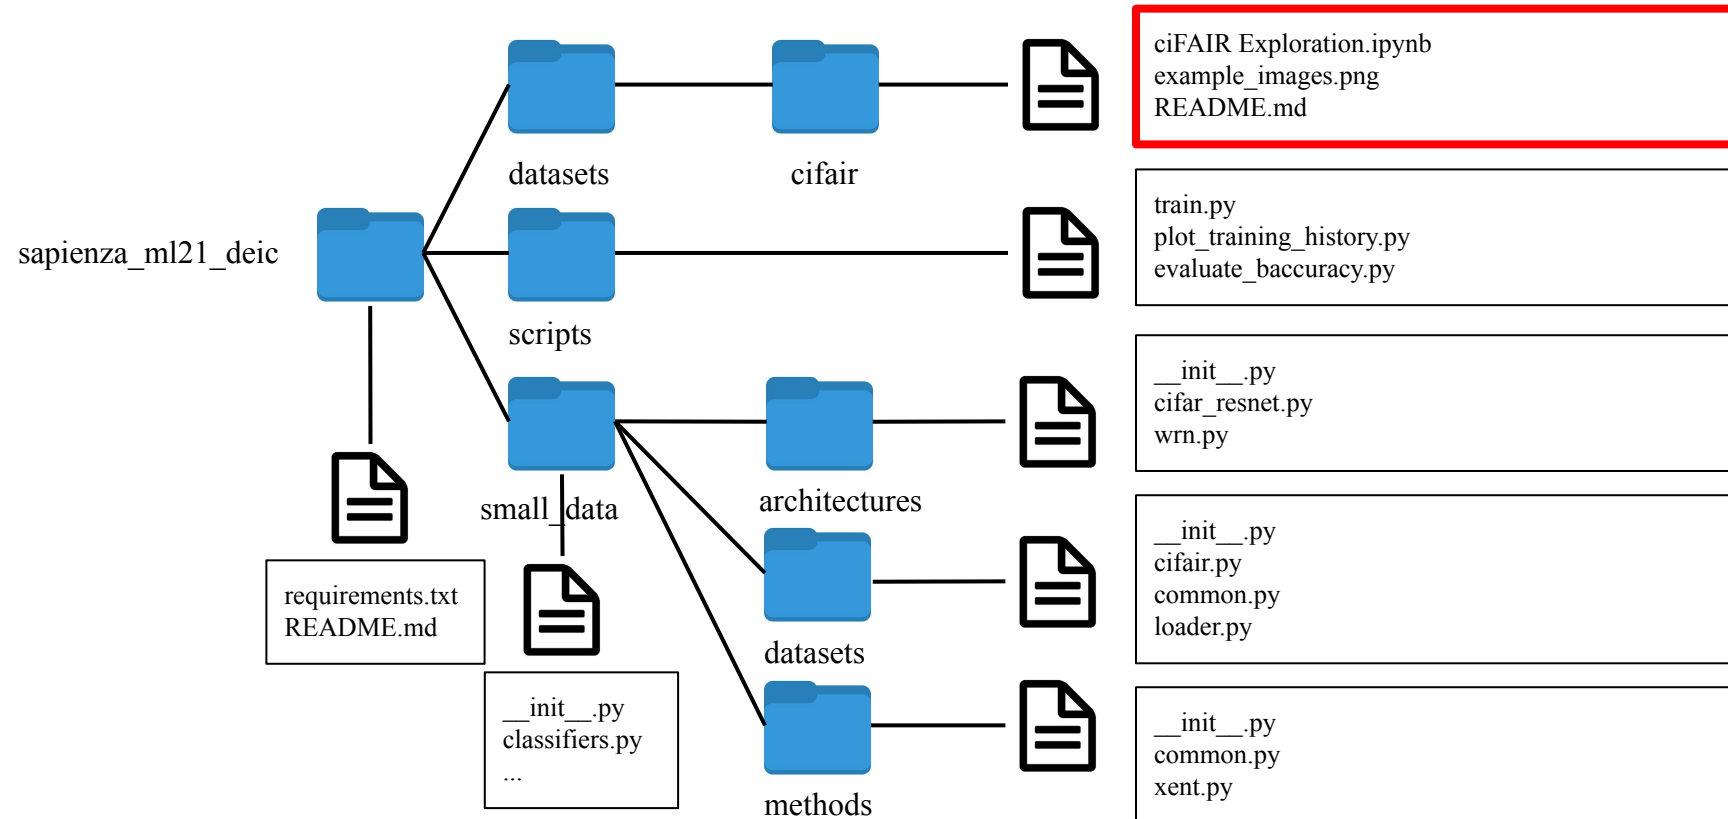
Put the provided dataset repository (ciFAIR-10 shared on Drive) inside this directory.





Code Repository

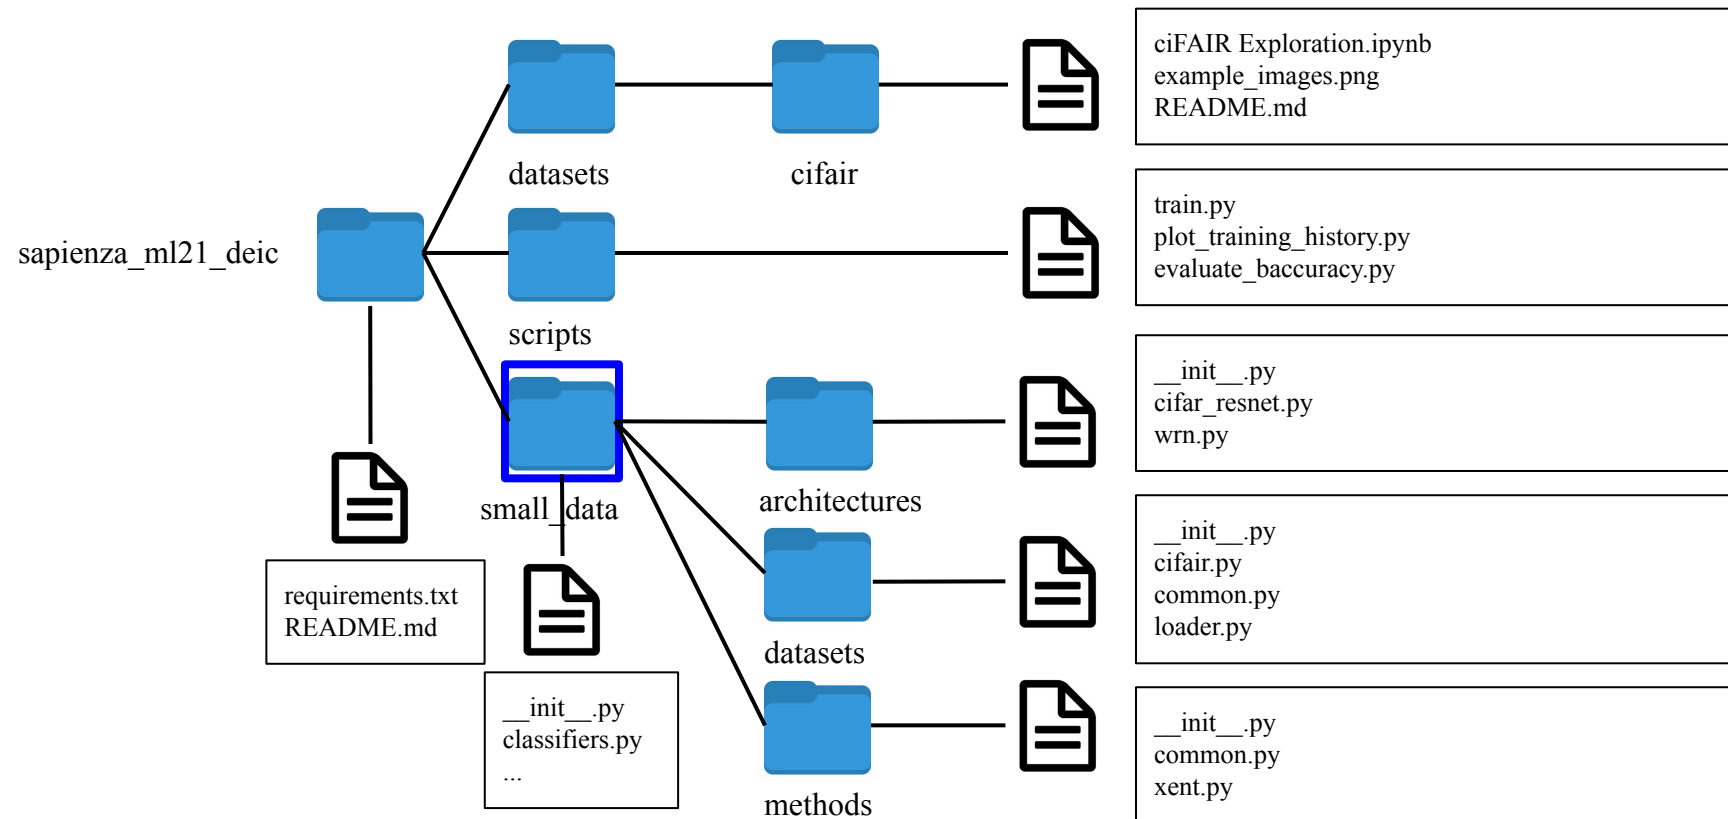
Notebook for dataset visualization and readme file with baseline performance along with terminal command to reproduce the baseline result.





Code Repository

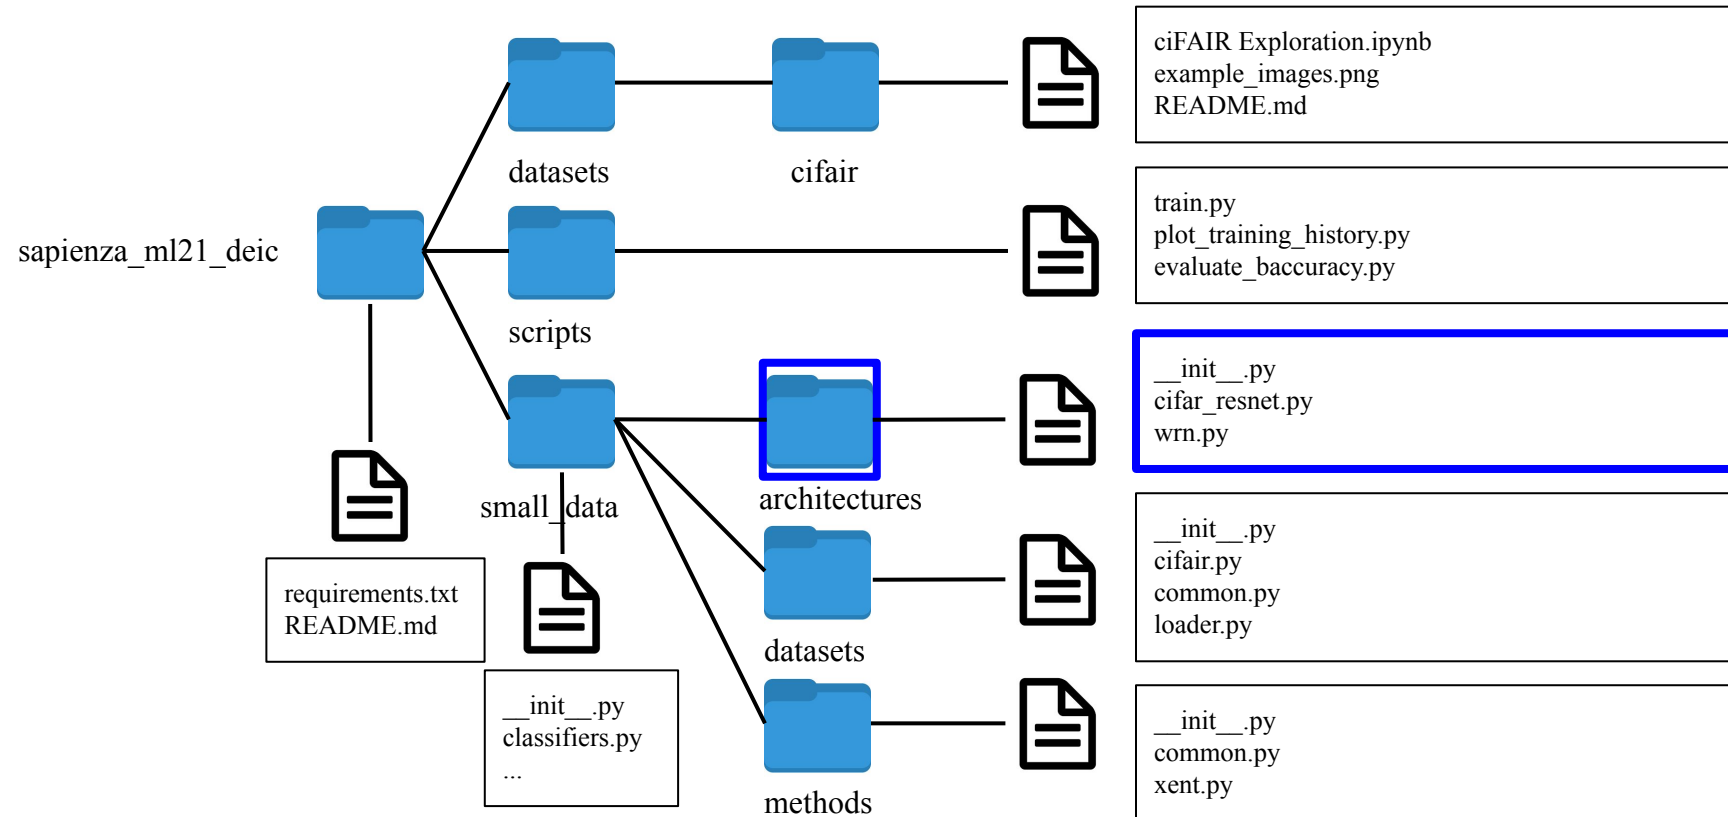
small_data is the directory that you will submit, with all its content.





Code Repository

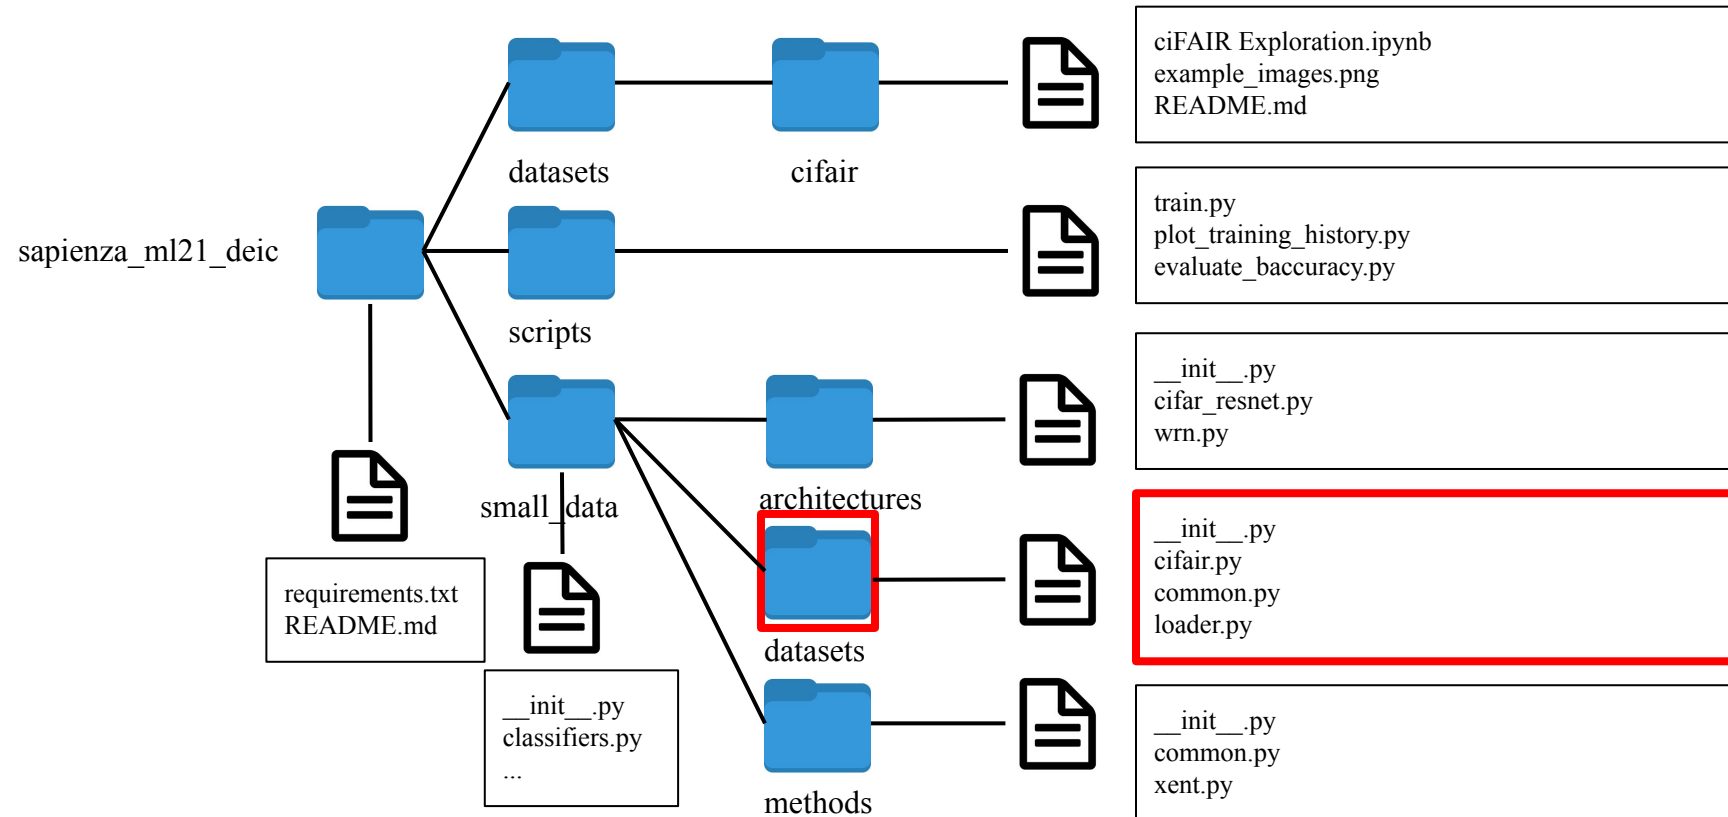
The folder **architectures** is the place to add other architecture files (e.g., densenet.py).





Code Repository

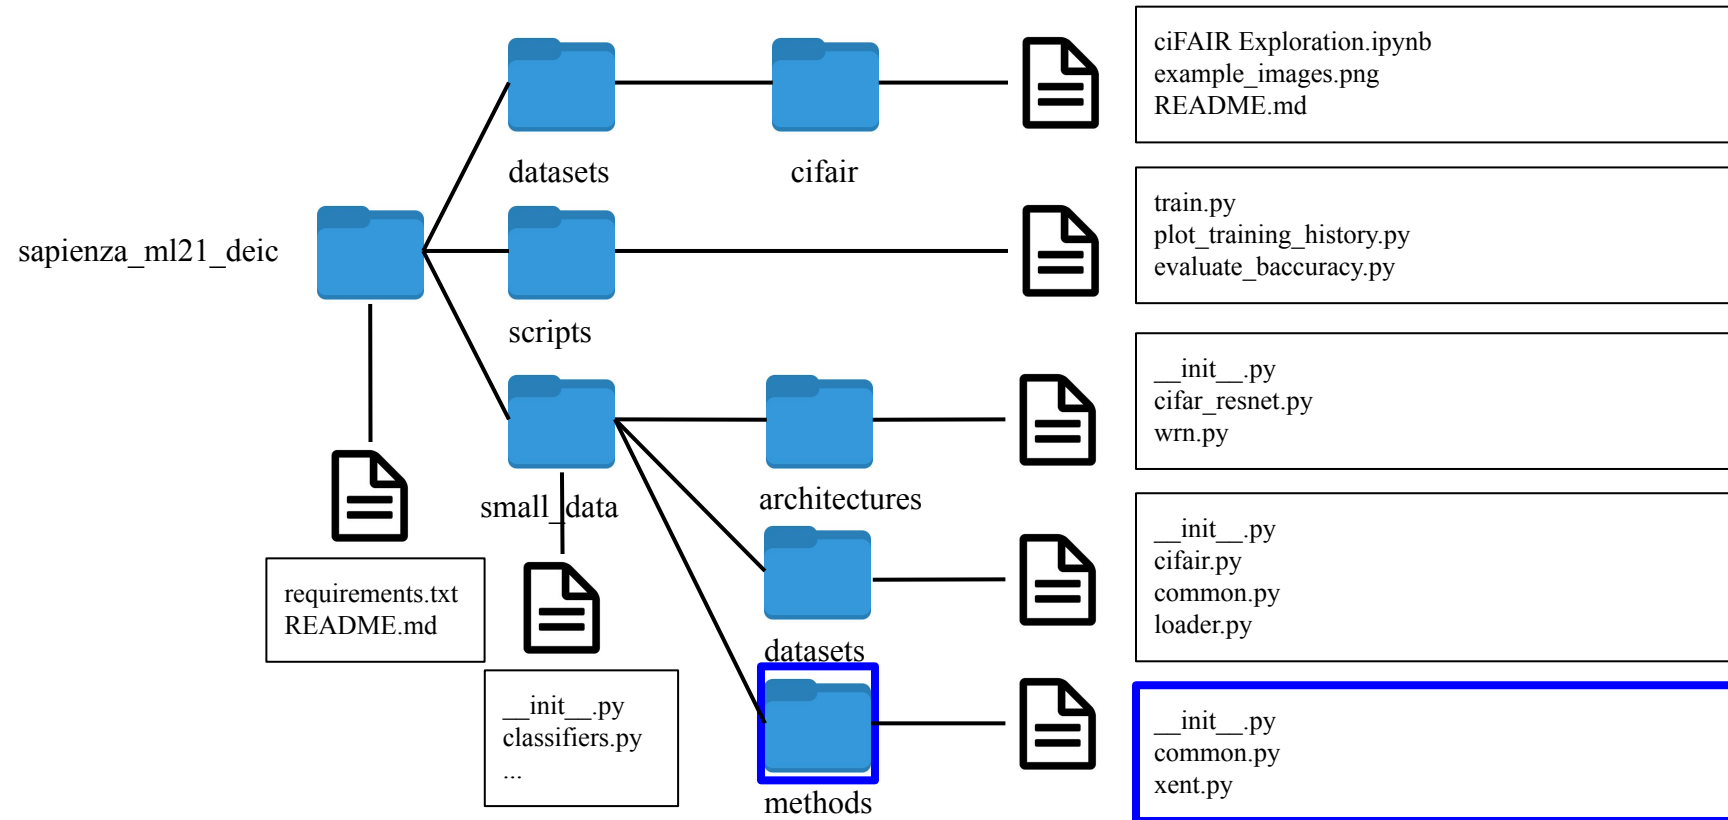
The folder **datasets** contains functions for loading data.





Code Repository

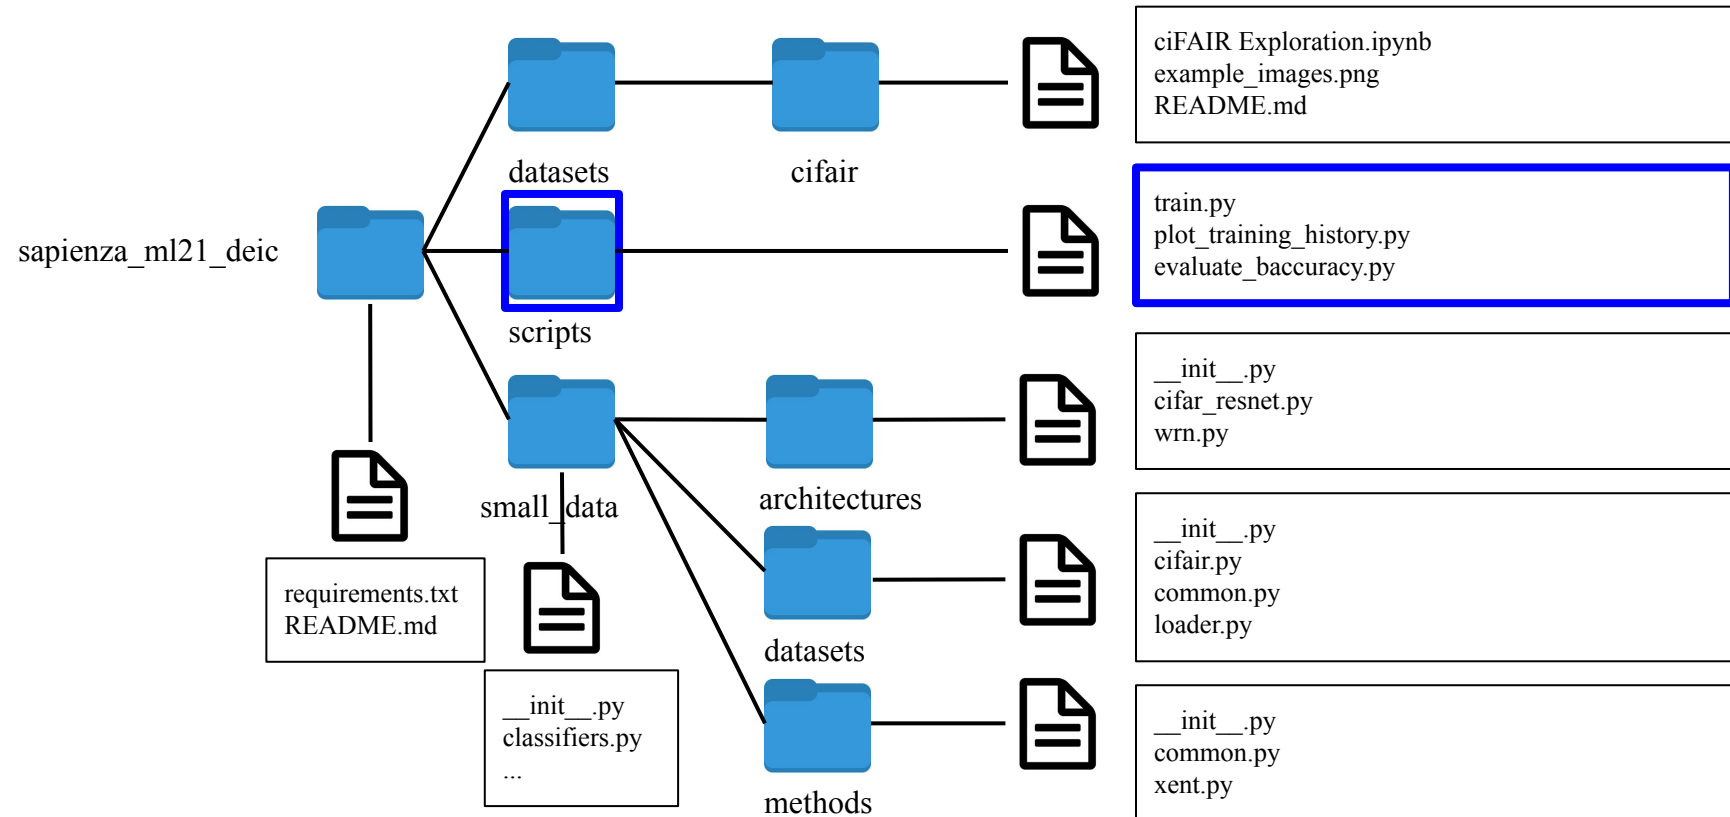
The folder **methods** contains functions for running the training algorithms.





Code Repository

The folder **scripts** contains higher-level scripts to actually run the training or visualize results.





SAPIENZA
UNIVERSITÀ DI ROMA

Google Colab Demo

- Setup the code repository on Google Drive with [SetupGitRepoOnDrive.ipynb](#)
- Run baseline experiment with [ExampleTrain.ipynb](#)

colab



Your Task


Add something in the code, save the trained model and evaluate it on the test set.

Examples:

- New architecture
 - A different optimizer for training
 - A choosable hyper-parameter (e.g., momentum parameter of SGD)
-
- **At least perform two of these tasks.**
 - **Only running the demo would translate in 0 points for your homework.**



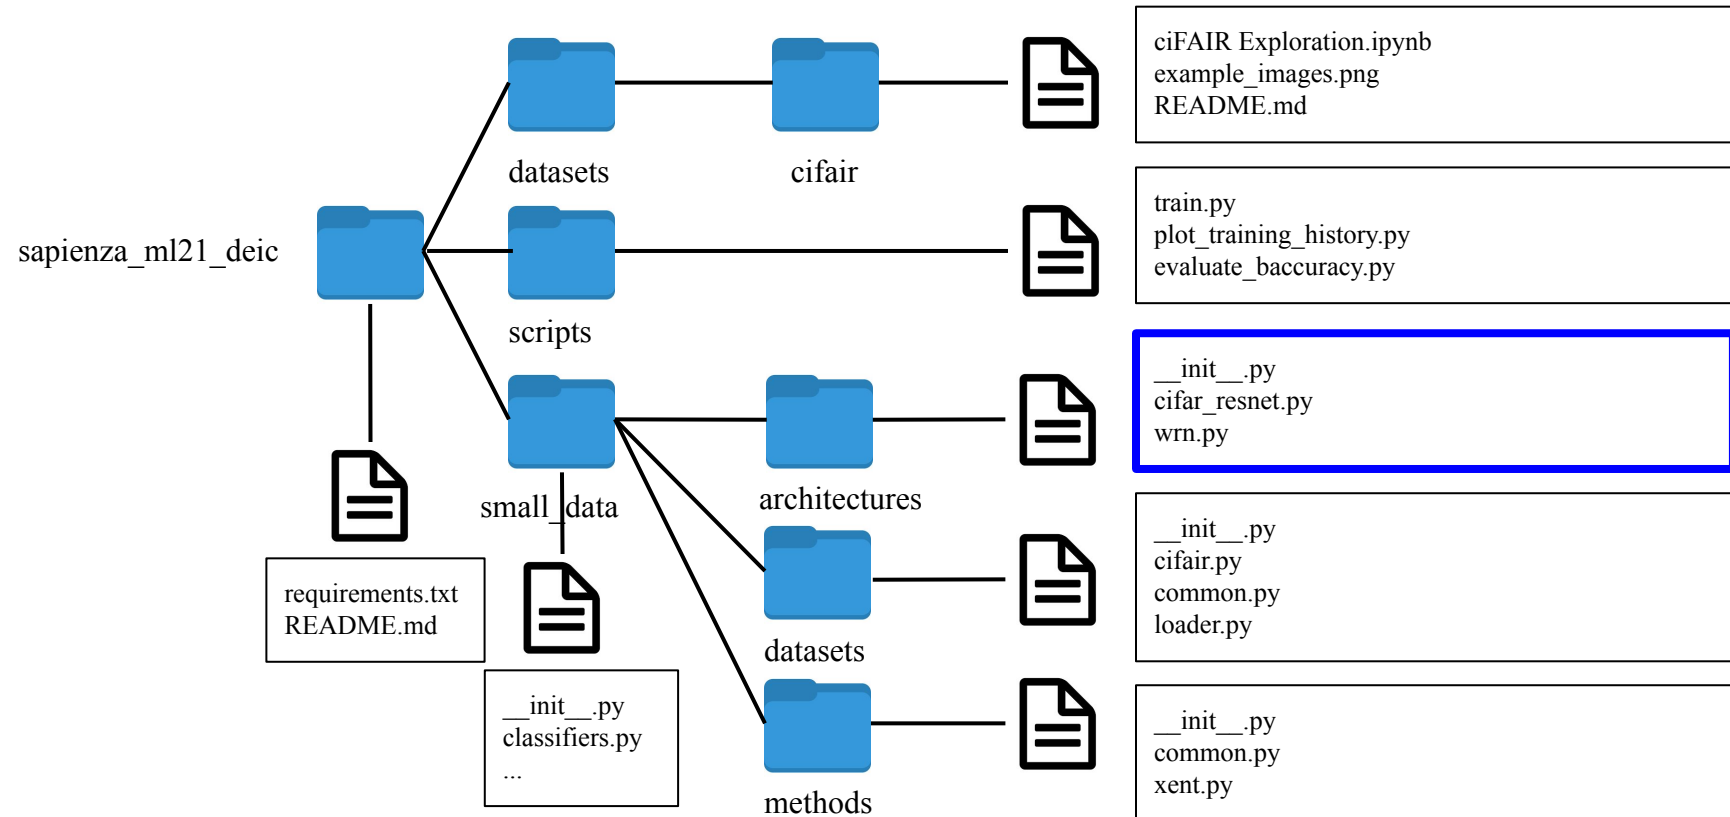
Adding a new architecture





Adding a new architecture

You should add a file that define the architecture inside **architectures**.





Adding a new architecture


Steps:

- Create a new file inside folder **architectures** (e.g., [wrn.py](#)) to contain all the python classes/functions needed by the given architecture.
- The main class must have an `__init__`, `forward`, `get_classifiers`, and `build_classifier` method.

```
class WideResNet(nn.Module):  
    def __init__(self, depth, num_classes, input_channels=3, widen_factor=1, ...):  
        ...  
  
    def forward(self, x):  
        ...  
  
    @staticmethod  
    def get_classifiers():  
        return ['wrn-16-8', 'wrn-16-10', ..., 'wrn-28-12']  
  
    @classmethod  
    def build_classifier(cls, arch: str, num_classes: int, input_channels: int):  
        _, depth, widen_factor = arch.split('-')  
        cls_instance = cls(int(depth), num_classes, input_channels=input_channels,  
                           widen_factor=int(widen_factor))  
        return cls_instance
```



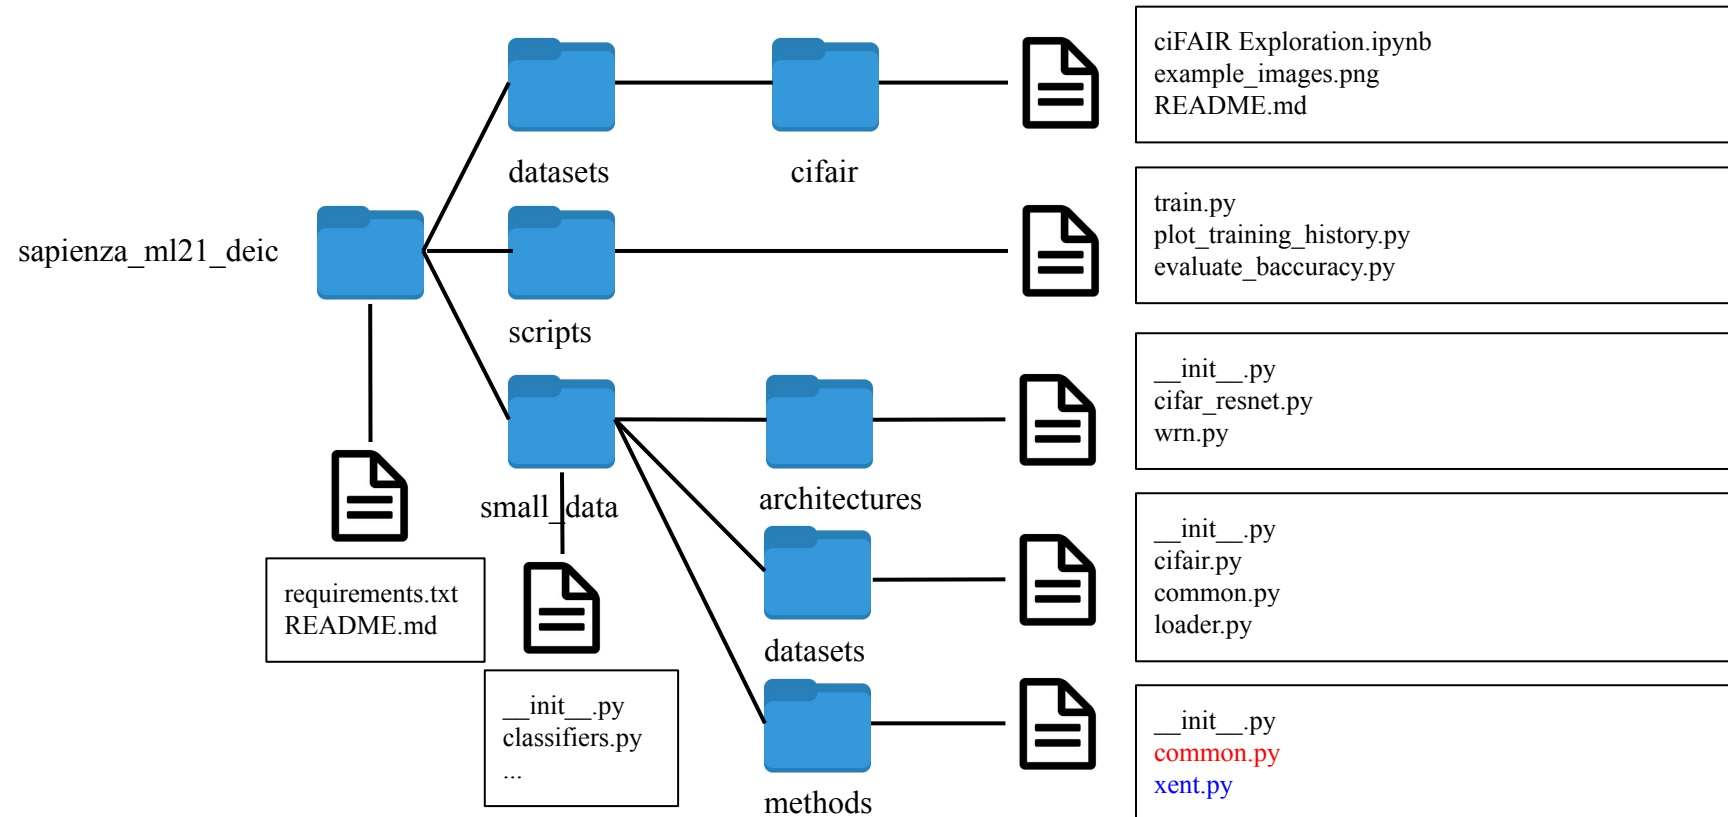
Using a different optimizer for training





Using a different optimizer for training

You should sub-class the method `get_optimizer` of `LearningMethod` defined in **common.py** with a different instance inside **xent.py**.





Using a different optimizer for training

The main training class contains multiple methods that can be sub-classed:

- `LearningMethod` is the general class for a training pipeline of a method defined in `common.py`.
- An instance of this class is `CrossEntropyClassifier` which is defined in `xent.py`.
- `CrossEntropyClassifier`, when called will override all methods of `LearningMethod` defined inside its file.
- Any change can be applied by modifying the desired class method (e.g., `get_optimizer`).



Using a different optimizer for training

Original `get_optimizer` function defined inside **common.py**.

```
def get_optimizer(self, model: nn.Module, max_epochs: int, max_iter: int):  
  
    optimizer = torch.optim.SGD(model.parameters(), lr=self.hparams['lr'], momentum=0.9,  
weight_decay=self.hparams['weight_decay'])  
  
    scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=max_iter)  
  
    return optimizer, scheduler
```



Using a different optimizer for training

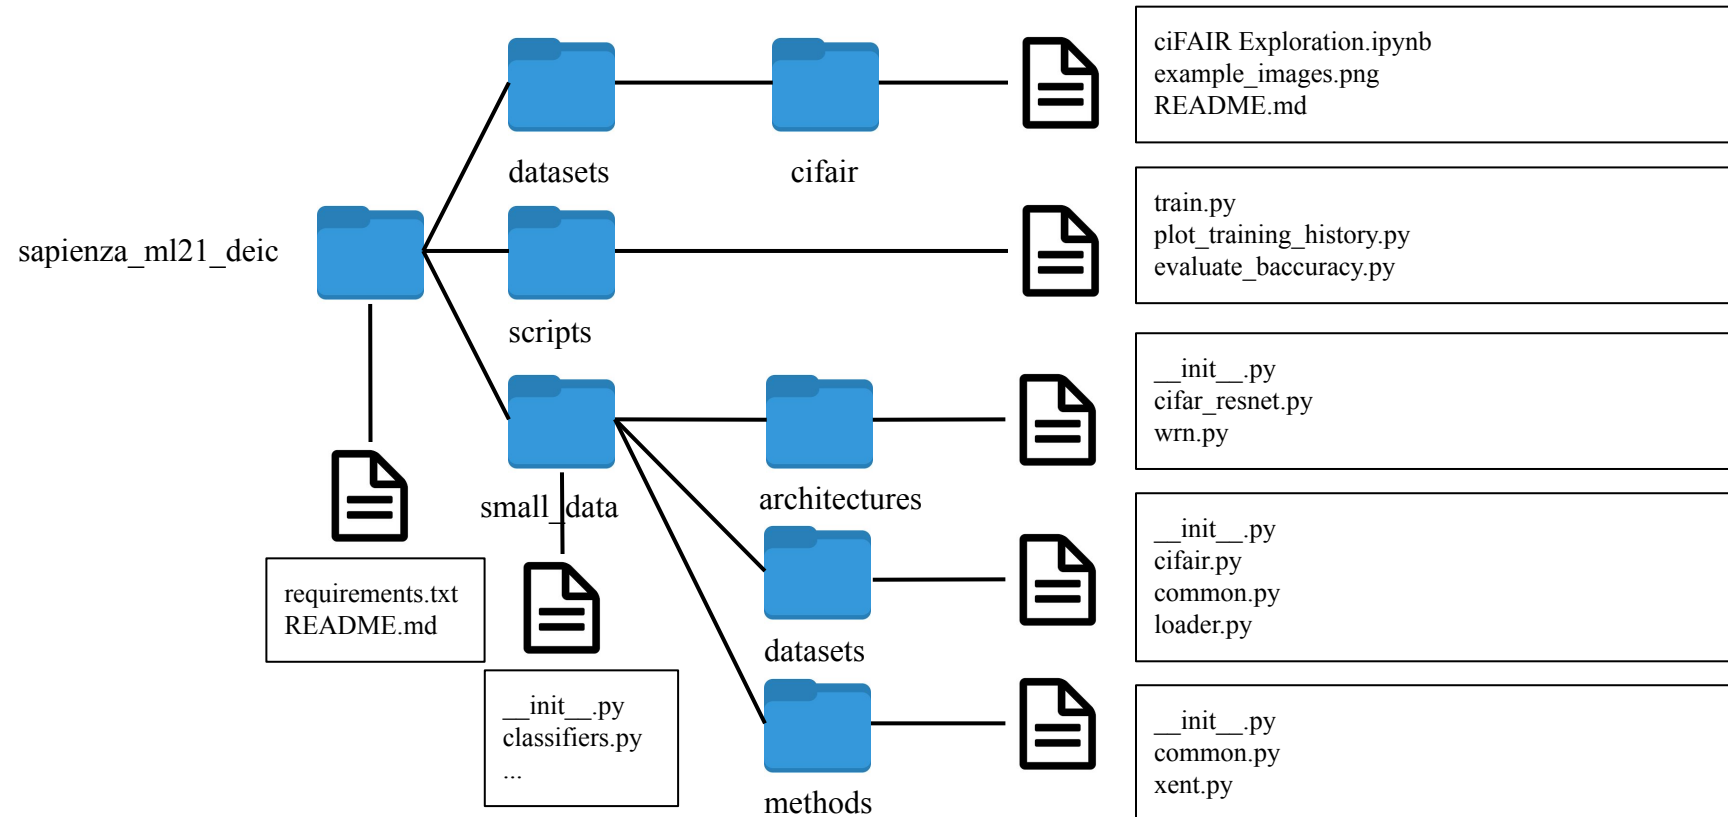
New `get_optimizer` method (with Adam instead of SGD) of `CrossEntropyClassifier` defined inside [xent.py](#).

```
def get_optimizer(self, model: nn.Module, max_epochs: int, max_iter: int):  
  
    optimizer = torch.optim.Adam(model.parameters(), lr=self.hparams['lr'],  
weight_decay=self.hparams['weight_decay'])  
  
    scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=max_iter)  
  
    return optimizer, scheduler
```




Add choosable hyper-parameter

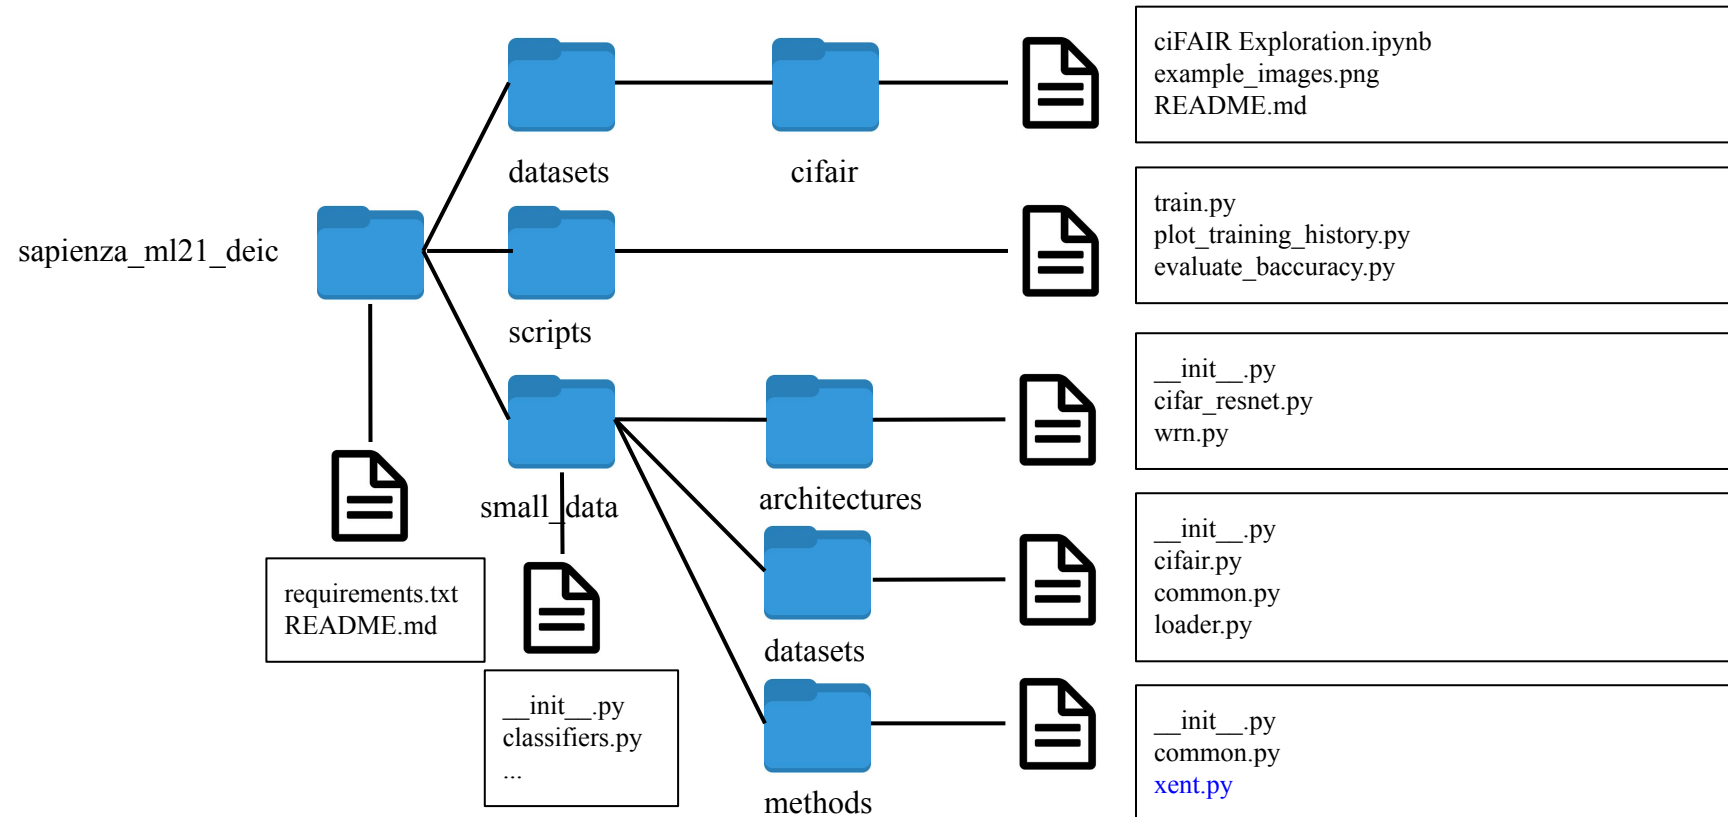
Making the momentum parameter of SGD as choosable hyper-parameter by the user.





Add choosable hyper-parameter

You should sub-class the static methods `default_hparams` and `get_optimizer` of class `CrossEntropyClassifier` in [xent.py](#).





Add choosable hyper-parameter

Sub-class the static `default_hparams` method of `CrossEntropyClassifier` to add momentum:

```
@staticmethod
def default_hparams() -> dict:
    return {
        **super(CrossEntropyClassifier, CrossEntropyClassifier).default_hparams(),
        'momentum' : 0.9
    }
```

Update the `get_optimizer` method too:

```
def get_optimizer(self, model: nn.Module, max_epochs: int, max_iter: int):

    optimizer = torch.optim.SGD(model.parameters(), lr=self.hparams['lr'],
    momentum=self.hparams['momentum'], weight_decay=self.hparams['weight_decay'])

    scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=max_iter)

    return optimizer, scheduler
```



Deep Learning Library

The code is written in PyTorch which is very similar to Tensorflow:

- PyTorch (1.7) documentation [[link](#)].
- The modules of `torch.nn` are equivalent to layers in `tf.keras.layers`. Docs [[link](#)].
- To use different network architectures I advise looking at the available models at `torchvision.models` [[link](#)] and copying, using/modifying their source code (e.g., EfficientNet).

EfficientNet

```
torchvision.models.efficientnet_b0(pretrained: bool = False, progress: bool = True,  
**kwargs: Any) → torchvision.models.efficientnet.EfficientNet [SOURCE]
```

```
class EfficientNet(nn.Module):  
    def __init__(  
        self,
```



Save the trained model

Models in PyTorch are saved similarly to Tensorflow:

- A `torch.nn.Module` has a state dictionary containing all the information concerning the layer weights.
- It is enough to simply call:

```
torch.save(model.state_dict(), "model.pth")
```
- Note that to load back the model, you should first instantiate an instance of the model class and then override the state dictionary.
- Therefore, you should call:

```
model.load_state_dict(torch.load("model.pth"))
```
- Additional documentation is available [[link](#)]



Suggestions

To start from a good baseline take into account these tricks:

- Small batch sizes (e.g., 10 or less) are preferable to improve generalization.
- An high number of epochs (e.g., > 500) allow the network to find better local minima with better generalization.
- Not all lr-wd couples work the same. Choosing a good couple strongly increases the results (Seminar 2) and may change with network width/depth.
- Network width ease the optimization process. With the right hyper-parameters, a wider network (i.e., more conv. filters) may find better solutions than thinner ones.
- Baseline to beat:
 - Wide ResNet-16-8 (wrn-16-8) , epochs = 500, bs = 10, lr = $4.55e-3$, wd = $5.29e-3$
 - Accuracy on test set: 58.22%
- **Note:** the homework evaluation won't be related to the test set performance obtained (unless for extremely bad results).



SAPIENZA
UNIVERSITÀ DI ROMA

Submission

Through Google Classroom:

- **PDF report** with description of the models and results on the test set.
- **ZIP file of small_data folder**, named with your matricola code (e.g, 1771597.zip).
- **Your trained model** saved with your matricola code (e.g, 1771597.pth).



Expected Outcome

Publish a paper with 100 names!! Or less optimistically:

- Get interesting insights and results about this crowd-sourcing experiment with ensembles on small datasets.
- Not getting the same model from all of you since this would basically ruin the experiment:
 - predictions of average of 100 identical models = predictions of 1 model
- And would give 0 points to your homework!!
- You'll learn important tools and practices to train neural networks.



Summary

To run the code you need a GPU either locally or on the Cloud:

- Use Colab or make sure to have PyTorch locally installed with GPU support
- [For Colab] Use the notebook [SetupGitRepoOnDrive.ipynb](#) to clone the repository on your Google Drive.
- [For Colab] Use/take inspiration from [ExampleTrain.ipynb](#) to run experiments.
- [For Colab] Note that to change the running code, you should modify the files of the repo that are located in your Google Drive.